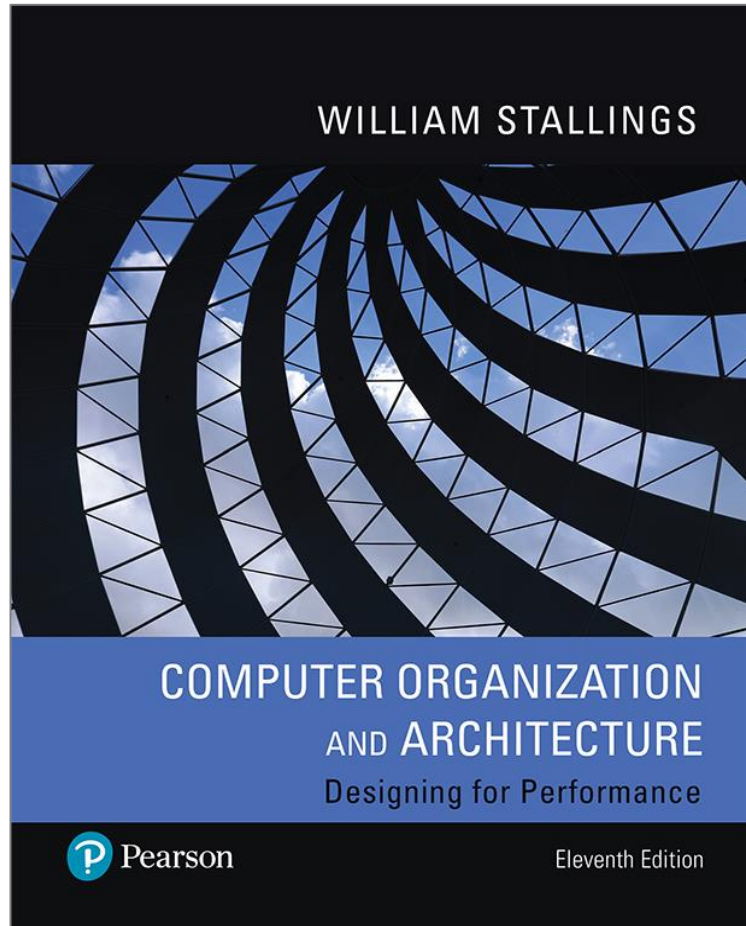


Computer Organization and Architecture

Designing for Performance

11th Edition



Chapter 8

Input/Output

Table 8.1

I/O Techniques

	No Interrupts	Use of Interrupts
I/O-to-memory transfer through processor	Programmed I/O	Interrupt-driven I/O
Direct I/O-to-memory transfer		Direct memory access (DMA)

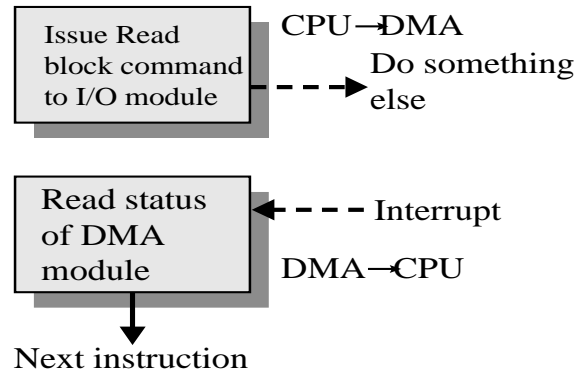
Drawbacks of Programmed and Interrupt-Driven I/O

- Both forms of I/O suffer from two inherent drawbacks:
 - 1) The I/O transfer rate is limited by the speed with which the processor can test and service a device
 - 2) The processor is tied up in managing an I/O transfer; a number of instructions must be executed for each I/O transfer
- There is a trade-off between these two drawbacks.
- Consider the transfer of a block of data. Using simple programmed I/O, the processor is dedicated to the task of I/O and can move data at a rather high rate, at the cost of doing nothing else.

Drawbacks of Programmed and Interrupt-Driven I/O

- Interrupt I/O frees up the processor to some extent at the expense of the I/O transfer rate. Nevertheless, both methods have an adverse impact on both processor activity and I/O transfer rate.
- When large volumes of data are to be moved a more efficient technique is direct memory access (DMA)

Figure 8.4 – Direct Memory Access (DMA)



(c) Direct memory access

- Direct memory access (DMA)
 - The I/O module and main memory exchange data directly without processor involvement

DMA

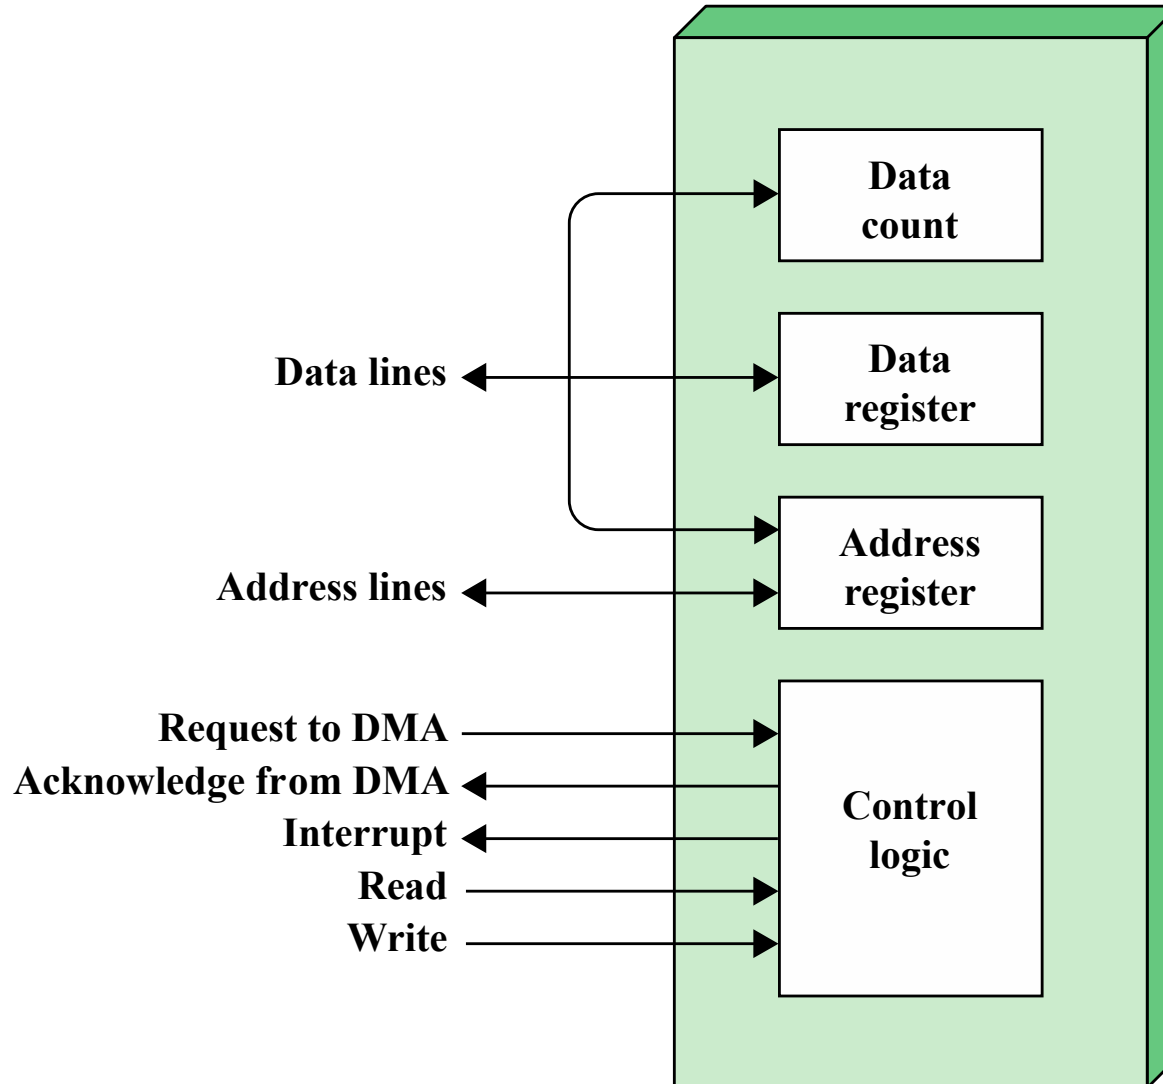
- DMA involves an additional module on the system bus.
- The DMA module is capable of mimicking the processor and, indeed, of taking over control of the system from the processor.
- It needs to do this to transfer data to and from memory over the system bus.

DMA

- For this purpose, the DMA module must use the bus only when the processor does not need it, or it must force the processor to suspend operation temporarily.
- The latter technique is more common and is referred to as cycle stealing, because the DMA module in effect steals a bus cycle.

Figure 8.12

Typical DMA Block Diagram



DMA


- When the processor wishes to read or write a block of data, it issues a command to the DMA module, by sending to the DMA module the following information:
 - Whether a read or write is requested, using the read or write control line between the processor and the DMA module
 - The address of the I/O device involved, communicated on the data lines
 - The starting location in memory to read from or write to, communicated on the data lines and stored by the DMA module in its address register
 - The number of words to be read or written, again communicated via the data lines and stored in the data count register

DMA


- The processor then continues with other work since it has **delegated this I/O operation to the DMA module.**
- The DMA module transfers the entire block of data, one word at a time, directly to or from memory, without going through the processor.
- When the transfer is complete, the DMA module sends an interrupt signal to the processor.
- Thus, the processor is involved only at the **beginning and end of the transfer.**

Figure 8.14

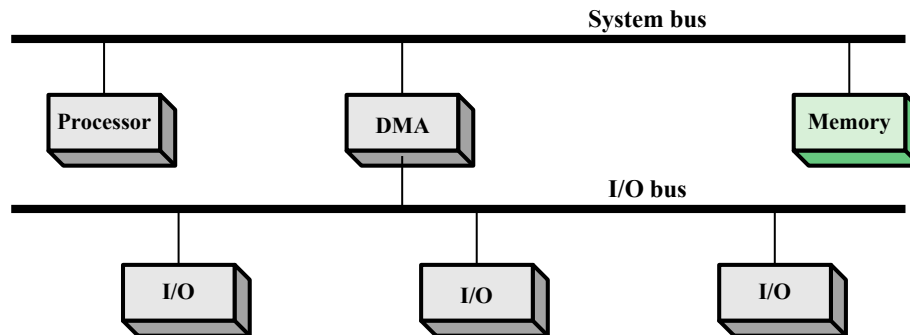
Alternative DMA Configurations



(a) Single-bus, detached DMA



(b) Single-bus, Integrated DMA-I/O



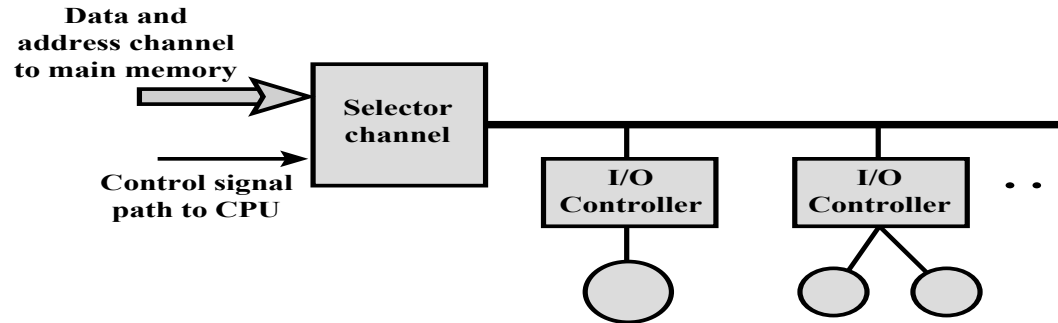
(c) I/O bus

Evolution of the I/O Function

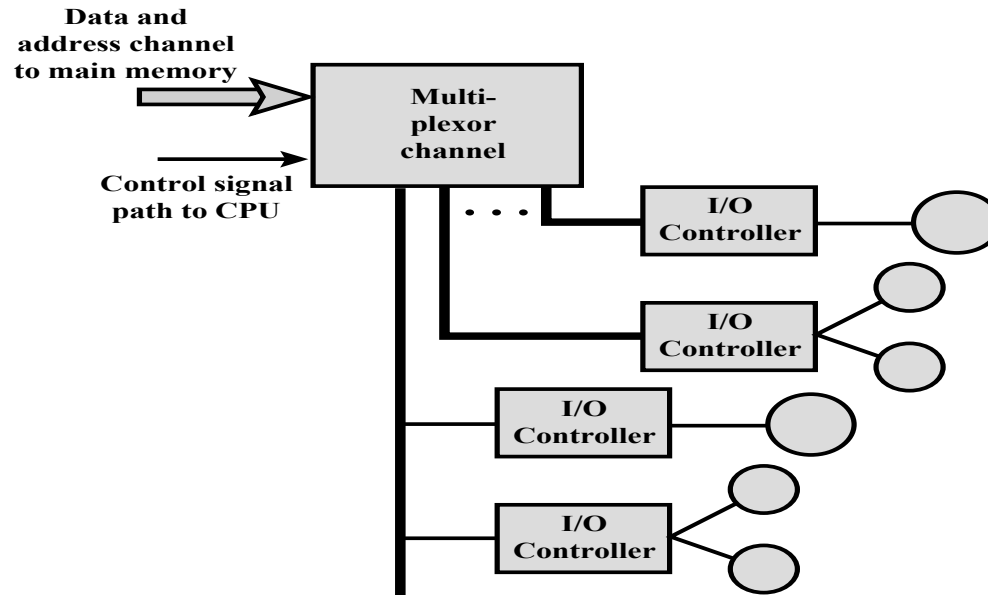
1. The CPU directly controls a peripheral device.
2. A controller or I/O module is added. The CPU uses programmed I/O without interrupts.
3. Same configuration as in step 2 is used, but now interrupts are employed. The CPU need not spend time waiting for an I/O operation to be performed, thus increasing efficiency.
4. The I/O module is given direct access to memory via DMA. It can now move a block of data to or from memory without involving the CPU, except at the beginning and end of the transfer.
5. The I/O module is enhanced to become a processor in its own right, with a specialized instruction set tailored for I/O
6. The I/O module has a local memory of its own and is, in fact, a computer in its own right. With this architecture a large set of I/O devices can be controlled with minimal CPU involvement.

Figure 8.18

I/O Channel Architecture



(a) Selector



(b) Multiplexor

I/O Channel

- A **selector channel** controls multiple high-speed devices and, at any one time, is dedicated to the transfer of data with one of those devices. Thus, the I/O channel selects one device and effects the data transfer.
- Each device, or a small set of devices, is handled by a controller, or I/O module, that is much like the I/O modules we have been discussing.
- Thus, the I/O channel serves in place of the CPU in controlling these I/O controllers.
- A **multiplexor channel** can handle I/O with multiple devices at the same time.

Copyright



This work is protected by United States copyright laws and is provided solely for the use of instructions in teaching their courses and assessing student learning. dissemination or sale of any part of this work (including on the World Wide Web) will destroy the integrity of the work and is not permitted. The work and materials from it should never be made available to students except by instructors using the accompanying text in their classes. All recipients of this work are expected to abide by these restrictions and to honor the intended pedagogical purposes and the needs of other instructors who rely on these materials.