


Chapter 2 :: Programming Language Syntax

Programming Language Pragmatics, Fourth Edition

Michael L. Scott



Regular Expressions

- A regular expression is one of the following:
 - A character
 - The empty string, denoted by ε
 - Two regular expressions concatenated
 - Two regular expressions separated by $|$ (i.e., or)
 - A regular expression followed by the Kleene star $*$ (concatenation of zero or more strings)

Regular Expressions

- Numerical constants accepted by a simple hand-held calculator:

number \longrightarrow *integer* | *real*

integer \longrightarrow *digit* *digit* *

real \longrightarrow *integer* *exponent* | *decimal* (*exponent* | ϵ)

decimal \longrightarrow *digit* * (. *digit* | *digit* .) *digit* *

exponent \longrightarrow (*e* | *E*) (+ | - | ϵ) *integer*

digit \longrightarrow 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9



Context-Free Grammars

- The notation for context-free grammars (CFG) is sometimes called Backus-Naur Form (BNF)
- A CFG consists of
 - A set of *terminals* T
 - A set of *non-terminals* N
 - A *start symbol* S (a non-terminal)
 - A set of *productions*



Context-Free Grammars

- Expression grammar with precedence and associativity

$$\text{expr} \longrightarrow \text{id} \mid \text{number} \mid - \text{expr} \mid (\text{expr}) \\ \mid \text{expr op expr}$$
$$\text{op} \longrightarrow + \mid - \mid * \mid /$$


Context-Free Grammars

- In this grammar,
generate the string

"slope * x + intercept"


$$\begin{aligned} \text{expr} &\longrightarrow \text{id} \mid \text{number} \mid - \text{expr} \mid (\text{expr}) \\ &\quad \mid \text{expr op expr} \\ \text{op} &\longrightarrow + \mid - \mid * \mid / \end{aligned}$$

$$\begin{aligned} \text{expr} &\Longrightarrow \text{expr op } \underline{\text{expr}} \\ &\Longrightarrow \text{expr } \underline{\text{op}} \text{ id} \\ &\Longrightarrow \underline{\text{expr}} + \text{id} \\ &\Longrightarrow \text{expr op } \underline{\text{expr}} + \text{id} \\ &\Longrightarrow \text{expr } \underline{\text{op}} \text{ id} + \text{id} \\ &\Longrightarrow \underline{\text{expr}} * \text{id} + \text{id} \\ &\Longrightarrow \text{id} * \text{id} + \text{id} \\ &\quad (\text{slope}) \quad (\text{x}) \quad (\text{intercept}) \end{aligned}$$




Context-Free Grammars

- Parse tree for expression grammar for **"slope * x + intercept"**



Context-Free Grammars

- Alternate (Incorrect) Parse tree for "slope * x + intercept"
- Our grammar is ambiguous



Context-Free Grammars

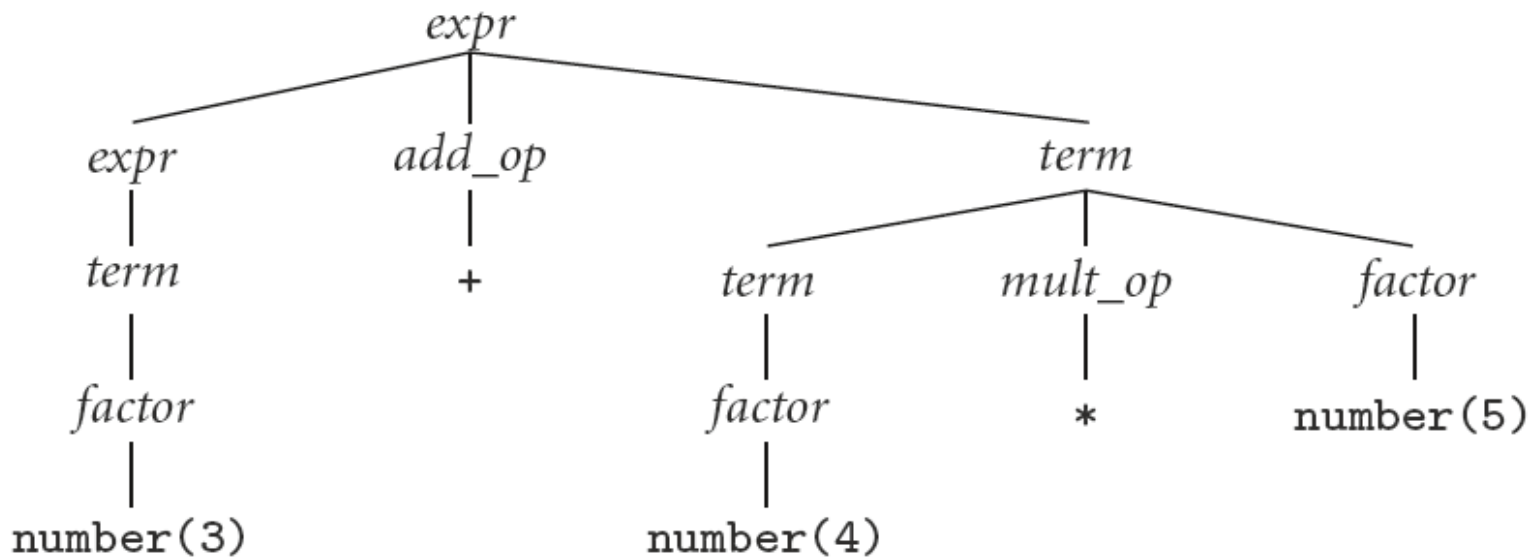
- A better version because it is unambiguous and captures precedence

1. $expr \longrightarrow term \mid expr \text{ add_op } term$
2. $term \longrightarrow factor \mid term \text{ mult_op } factor$
3. $factor \longrightarrow id \mid number \mid - factor \mid (expr)$
4. $add_op \longrightarrow + \mid -$
5. $mult_op \longrightarrow * \mid /$



Context-Free Grammars

- Parse tree for expression grammar (with left associativity) for **3 + 4 * 5**



Scanning

- Recall scanner is responsible for
 - tokenizing source
 - removing comments
 - (often) dealing with *pragmas* (i.e., significant comments)
 - saving text of identifiers, numbers, strings
 - saving source locations (file, line, column) for error messages

Scanning

- Suppose we are building an ad-hoc (hand-written) scanner for Pascal:
 - We read the characters one at a time with look-ahead
- If it is one of the one-character tokens
`{ () [] < > , ; = + - etc }`
we announce that token
- If it is a `.`, we look at the next character
 - If that is a dot, we announce `.`
 - Otherwise, we announce `.` and reuse the look-ahead

Scanning

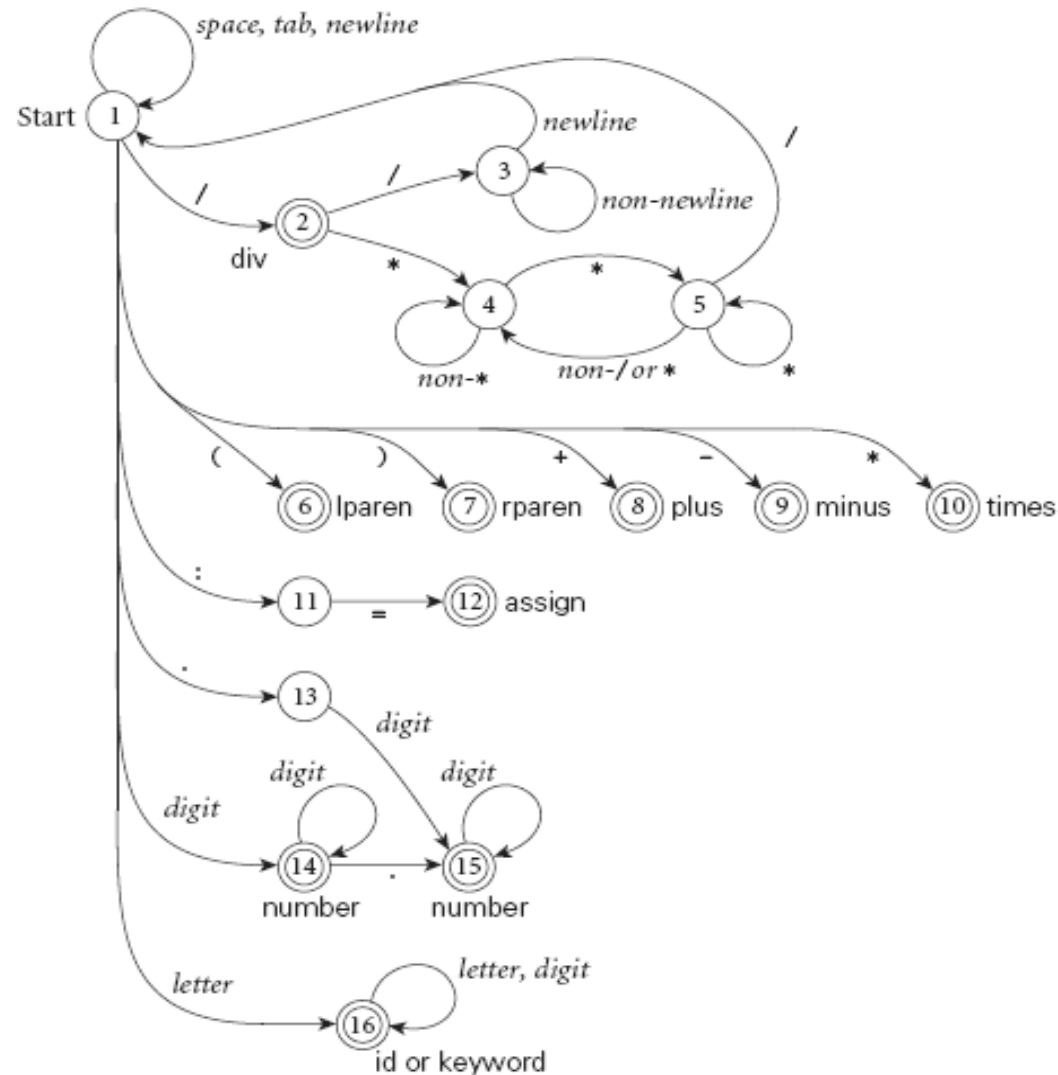
- If it is a $<$, we look at the next character
 - if that is a $=$ we announce $<=$
 - otherwise, we announce $<$ and reuse the look-ahead, etc
- If it is a letter, we keep reading letters and digits and maybe underscores until we can't anymore
 - then we check to see if it is a reserve word

Scanning

- If it is a digit, we keep reading until we find a non-digit
 - if that is not a . we announce an integer
 - otherwise, we keep looking for a real number
 - if the character after the . is not a digit we announce an integer and reuse the . and the look-ahead

Scanning

- Pictorial representation of a scanner for calculator tokens, in the form of a finite automaton



Scanning

- This is a deterministic finite automaton (DFA)
 - Lex, scangen, etc. build these things automatically from a set of regular expressions
 - Specifically, they construct a machine that accepts the language

```
identifier | int const
| real const | comment | symbol
| ...
```


Scanning

- We run the machine over and over to get one token after another
 - Nearly universal rule:
 - always take the longest possible token from the input
thus foobar is foobar and never f or foo or foob
 - more to the point, 3.14159 is a real const and never 3, ., and 14159
- Regular expressions "generate" a regular language; DFAs "recognize" it

Scanning

- Scanners tend to be built three ways
 - ad-hoc
 - semi-mechanical pure DFA
(usually realized as nested case statements)
 - table-driven DFA
- Ad-hoc generally yields the fastest, most compact code by doing lots of special-purpose things, though good automatically-generated scanners come very close

Scanning

- Writing a pure DFA as a set of nested case statements is a surprisingly useful programming technique
 - though it's often easier to use perl, awk, sed
 - for details see Figure 2.11
- Table-driven DFA is what lex and scangen produce
 - lex (flex) in the form of C code
 - scangen in the form of numeric tables and a separate driver (for details see Figure 2.12)

Scanning

- Note that the rule about longest-possible tokens means you return only when the next character can't be used to continue the current token
 - the next character will generally need to be saved for the next token
- In some cases, you may need to peek at more than one character of look-ahead in order to know whether to proceed
 - In Pascal, for example, when you have a 3 and you see a dot
 - do you proceed (in hopes of getting 3.14)?
 - or
 - do you stop (in fear of getting 3..5)?

Scanning

- In messier cases, you may not be able to get by with any fixed amount of look-ahead. In Fortran, for example, we have
DO 5 I = 1, 25 loop
DO 5 I = 1.25 assignment
- Here, we need to remember we were in a potentially final state, and save enough information that we can back up to it, if we get stuck later

Parsing

- Terminology:
 - context-free grammar (CFG)
 - symbols
 - terminals (tokens)
 - non-terminals
 - production
 - derivations (left-most and right-most - canonical)
 - parse trees
 - sentential form

Parsing

- By analogy to RE and DFAs, a context-free grammar (CFG) is a *generator* for a context-free language (CFL)
 - a parser is a language *recognizer*
- There is an infinite number of grammars for every context-free language
 - not all grammars are created equal, however

Parsing

- It turns out that for any CFG we can create a parser that runs in $O(n^3)$ time
- There are two well-known parsing algorithms that permit this
 - Early's algorithm
 - Cooke-Younger-Kasami (CYK) algorithm
- $O(n^3)$ time is clearly unacceptable for a parser in a compiler - too slow

Parsing

- Fortunately, there are large classes of grammars for which we can build parsers that run in linear time
 - The two most important classes are called **LL** and **LR**
- LL stands for 'Left-to-right, Leftmost derivation'.
- LR stands for 'Left-to-right, Rightmost derivation'

Parsing

- LL parsers are also called 'top-down', or 'predictive' parsers & LR parsers are also called 'bottom-up', or 'shift-reduce' parsers
- There are several important sub-classes of LR parsers
 - SLR
 - LALR
- We won't be going into detail on the differences between them

Parsing

- Every LL(1) grammar is also LR(1), though right recursion in production tends to require very deep stacks and complicates semantic analysis
- Every CFL that can be parsed deterministically has an SLR(1) grammar (which is LR(1))
- Every deterministic CFL with the *prefix property* (no valid string is a prefix of another valid string) has an LR(0) grammar

Parsing

- You commonly see LL or LR (or whatever) written with a number in parentheses after it
 - This number indicates how many tokens of look-ahead are required in order to parse
 - Almost all real compilers use one token of look-ahead
- The expression grammar (with precedence and associativity) you saw before is LR(1), but not LL(1)

LL Parsing

- Here is an LL(1) grammar (Fig 2.15):

```
1. program      → stmt list $$$
2. stmt_list    → stmt stmt_list
3.              | ε
4. stmt         → id := expr
5.              | read id
6.              | write expr
7. expr         → term term_tail
8. term_tail    → add op term term_tail
9.              | ε
```

LL Parsing

- LL(1) grammar (continued)

10. `term` \rightarrow `factor fact_tail`

11. `fact_tail` \rightarrow `mult_op fact fact_tail`

- $\mid \epsilon$
- `factor` \rightarrow `(expr)`
- \mid `id`
- \mid `number`
- `add_op` \rightarrow `+`
- \mid `-`
- `mult_op` \rightarrow `*`
- \mid `/`

LL Parsing

- Like the bottom-up grammar, this one captures associativity and precedence, but most people don't find it as pretty
 - for one thing, the operands of a given operator aren't in a RHS together!
 - however, the simplicity of the parsing algorithm makes up for this weakness
- How do we parse a string with this grammar?
 - by building the parse tree incrementally

LL Parsing

- Example (average program)

```
read A
```

```
read B
```

```
sum := A + B
```

```
write sum
```

```
write sum / 2
```

- We start at the top and predict needed productions on the basis of the current left-most non-terminal in the tree and the current input token

- Parse tree for the average program (Figure 2.18)



LL Parsing

- Table-driven LL parsing: you have a big loop in which you repeatedly look up an action in a two-dimensional table based on current leftmost non-terminal and current input token. The actions are
 - (1) match a terminal
 - (2) predict a production
 - (3) announce a syntax error

LL Parsing

- LL(1) parse table for parsing for calculator language

Top-of-stack nonterminal	Current input token											
	id	number	read	write	:=	()	+	-	*	/	\$\$
<i>program</i>	1	—	1	1	—	—	—	—	—	—	—	1
<i>stmt_list</i>	2	—	2	2	—	—	—	—	—	—	—	3
<i>stmt</i>	4	—	5	6	—	—	—	—	—	—	—	—
<i>expr</i>	7	7	—	—	—	7	—	—	—	—	—	—
<i>term_tail</i>	9	—	9	9	—	—	9	8	8	—	—	9
<i>term</i>	10	10	—	—	—	10	—	—	—	—	—	—
<i>factor_tail</i>	12	—	12	12	—	—	12	12	12	11	11	12
<i>factor</i>	14	15	—	—	—	13	—	—	—	—	—	—
<i>add_op</i>	—	—	—	—	—	—	—	16	17	—	—	—
<i>mult_op</i>	—	—	—	—	—	—	—	—	—	18	19	—



LL Parsing

- To keep track of the left-most non-terminal, you push the as-yet-unseen portions of productions onto a stack
 - for details see Figure 2.21
- The key thing to keep in mind is that the stack contains all the stuff you expect to see between now and the end of the program
 - what you *predict* you will see

LL Parsing

- Problems trying to make a grammar LL(1)
 - left recursion

- example:

`id_list → id | id_list , id`

equivalently

`id_list → id id_list_tail`

`id_list_tail → , id id_list_tail
 | epsilon`

- we can get rid of all left recursion mechanically in any grammar

LL Parsing

- Problems trying to make a grammar LL(1)
 - common prefixes: another thing that LL parsers can't handle
 - solved by "left-factoring"
 - example:
$$\text{stmt} \rightarrow \text{id} := \text{expr} \mid \text{id} (\text{arg_list})$$

equivalently
$$\text{stmt} \rightarrow \text{id id_stmt_tail}$$

$$\text{id_stmt_tail} \rightarrow := \text{expr}$$

$$\mid (\text{arg_list})$$
- we can eliminate left-factor mechanically



LL Parsing

- Note that eliminating left recursion and common prefixes does NOT make a grammar LL
 - there are infinitely many non-LL LANGUAGES, and the mechanical transformations work on them just fine
 - the few that arise in practice, however, can generally be handled with kludges

LL Parsing

- Problems trying to make a grammar LL(1)
 - the "dangling else" problem prevents grammars from being LL(1) (or in fact LL(k) for any k)
 - the following natural grammar fragment is ambiguous (Pascal)

```
stmt → if cond then_clause else_clause
      | other_stuff
then_clause → then stmt
else_clause → else stmt
             | epsilon
```


LL Parsing

- The less natural grammar fragment can be parsed bottom-up but not top-down

```
stmt → balanced stmt | unbalanced stmt
```

```
balanced_stmt → if cond then balanced_stmt
                else balanced_stmt
                | other stuff
```

```
unbalanced_stmt → if cond then stmt
                  | if cond then balanced_stmt
                  else      unbalanced_stmt
```



LL Parsing

- The usual approach, whether top-down OR bottom-up, is to use the ambiguous grammar together with a *disambiguating rule* that says
 - else goes with the closest then or
 - more generally, the first of two possible productions is the one to predict (or reduce)

LL Parsing

- Better yet, languages (since Pascal) generally employ explicit end-markers, which eliminate this problem
- In Modula-2, for example, one says:

```
if A = B then
    if C = D then E := F end
else
    G := H
end
```

- Ada says 'end if'; other languages say 'fi'



LL Parsing

- One problem with end markers is that they tend to bunch up. In Pascal you say

```
if A = B then ...  
else if A = C then ...  
else if A = D then ...  
else if A = E then ...  
else ...;
```

- With end markers this becomes

```
if A = B then ...  
else if A = C then ...  
else if A = D then ...  
else if A = E then ...  
else ...;  
end; end; end; end;
```

LL Parsing

- The algorithm to build predict sets is tedious (for a "real" sized grammar), but relatively simple
- It consists of three stages:
 - (1) compute FIRST sets for symbols
 - (2) compute FOLLOW sets for non-terminals (this requires computing FIRST sets for some *strings*)
 - (3) compute predict sets or table for all productions

LL Parsing

- Algorithm First/Follow/Predict:

- $\text{FIRST}(\alpha) == \{a : \alpha \rightarrow^* a \beta\}$
 $\cup (\text{if } \alpha \Rightarrow^* \varepsilon \text{ THEN } \{\varepsilon\} \text{ ELSE NULL})$
- $\text{FOLLOW}(A) == \{a : S \rightarrow^+ \alpha A a \beta\}$
 $\cup (\text{if } S \rightarrow^* \alpha A \text{ THEN } \{\varepsilon\} \text{ ELSE NULL})$
- $\text{Predict}(A \rightarrow X_1 \dots X_m) == (\text{FIRST}(X_1 \dots X_m) - \{\varepsilon\}) \cup (\text{if } X_1, \dots, X_m \rightarrow^* \varepsilon \text{ then FOLLOW}(A) \text{ ELSE NULL})$

- For examples, look at lecture notes.

LL Parsing

- If any token belongs to the predict set of more than one production with the same LHS, then the grammar is not LL(1)

LR Parsing

- LR parsers are almost always table-driven:
 - like a table-driven LL parser, an LR parser uses a big loop in which it repeatedly inspects a two-dimensional table to find out what action to take
 - unlike the LL parser, however, the LR driver has non-trivial state (like a DFA), and the table is indexed by current input token and current state
 - the stack contains a record of what has been seen SO FAR (NOT what is expected)

LR Parsing

- A scanner is a DFA
 - it can be specified with a state diagram
- An LL or LR parser is a PDA
 - Early's & CYK algorithms do NOT use PDAs
 - a PDA can be specified with a state diagram and a stack
 - the state diagram looks just like a DFA state diagram, except the arcs are labeled with <input symbol, top-of-stack symbol> pairs, and in addition to moving to a new state the PDA has the option of pushing or popping a finite number of symbols onto/off the stack



LR Parsing

- An SLR/LALR/LR PDA has multiple states
 - it is a "recognizer," not a "predictor"
 - it builds a parse tree from the bottom up
 - the states keep track of which productions we *might* be in the middle
- The parsing of the Characteristic Finite State Machine (CFSM) is based on
 - Shift
 - Reduce

LR Parsing

- This grammar is SLR(1), a particularly nice class of bottom-up grammar
 - it isn't exactly what we saw originally
 - we've eliminated the epsilon production to simplify the presentation
- For an example, look at Bottom-Up parsers lecture notes.